

HW 2: Time Command

Task 1. Implement a time1 command that reports elapsed time.

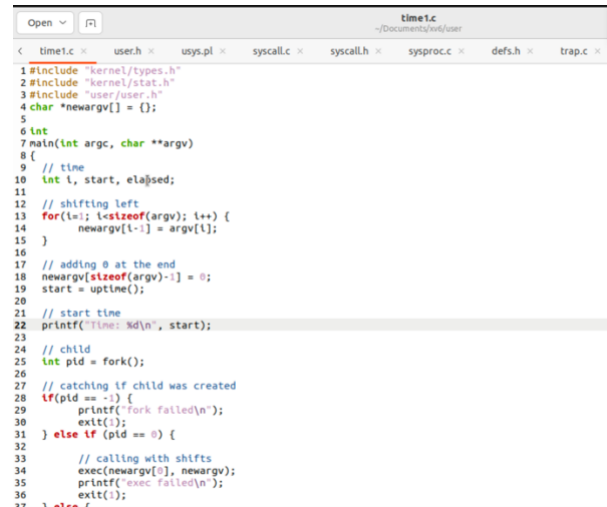
Figure 1: the code implemented that used the uptime to record the elapsed time of matmul.

Figure 2: Results from running my time1 command.

Figure 3: The time1, matmul, and sleep in UPROGS in Makefile.

I learned how to create a xv6 command, that can calculate the execution time elapse. I learned the correct integration into the build process, and now I comprehend the difference between working in the kernel or user.

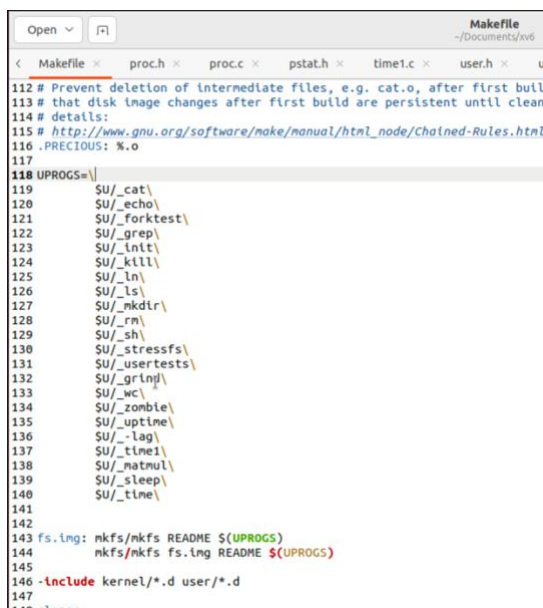
The difficulties I had were debugging problems in my custom time1 command.



```
time1.c
~/Documents/xv6/user

1 #include "kernel/types.h"
2 #include "kernel/stat.h"
3 #include "user/user.h"
4 char *newargv[] = {};
5
6 int
7 main(int argc, char **argv)
8 {
9     // time
10    int t, start, elapsed;
11
12    // shifting left
13    for (i=1; i<sizeof(argv); i++) {
14        newargv[i-1] = argv[i];
15    }
16
17    // adding 0 at the end
18    newargv[sizeof(argv)-1] = 0;
19    start = uptime();
20
21    // start time
22    printf("Time: %d\n", start);
23
24    // child
25    int pid = fork();
26
27    // catching if child was created
28    if (pid == -1) {
29        printf("fork failed\n");
30        exit(1);
31    } else if (pid == 0) {
32        // calling with shifts
33        exec(newargv[0], newargv);
34        printf("exec failed\n");
35        exit(1);
36    } else {
37        // ...
38    }
39 }
```

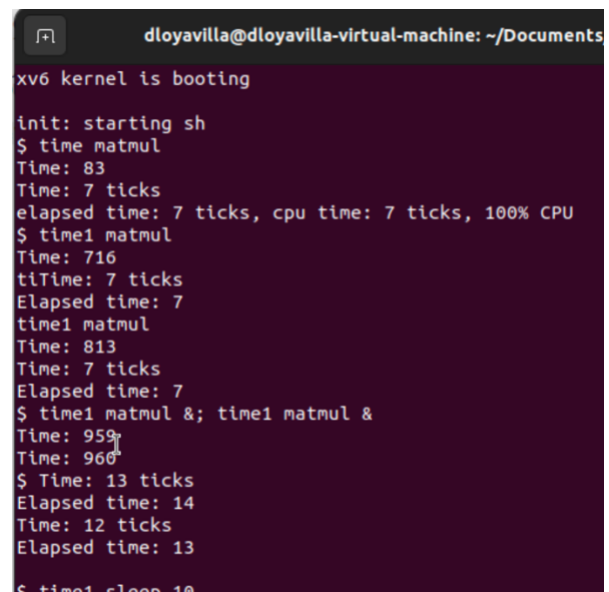
Figure 1. Time 1 Code



```
Makefile
~/Documents/xv6

112 # Prevent deletion of intermediate files, e.g. cat.o, after first build
113 # that disk image changes after first build are persistent until clean
114 # details:
115 # http://www.gnu.org/software/make/manual/html_node/Chained-Rules.html
116 .PRECIOUS: %.o
117
118 UPROGS=\
119     $U/_cat\
120     $U/_echo\
121     $U/_forktest\
122     $U/_grep\
123     $U/_init\
124     $U/_kill\
125     $U/_ln\
126     $U/_ls\
127     $U/_mkdir\
128     $U/_rm\
129     $U/_sh\
130     $U/_stressfs\
131     $U/_userstests\
132     $U/_grind\
133     $U/_wc\
134     $U/_zombie\
135     $U/_uptime\
136     $U/_lag\
137     $U/_time1\
138     $U/_matmul\
139     $U/_sleep\
140     $U/_time\
141
142
143 fs.img: mkfs/mkfs README $(UPROGS)
144     mkfs/mkfs fs.img README $(UPROGS)
145
146 -include kernel/*.d user/*.d
147
148 ...
```

Figure 4. UPROGS changes



```
dloyavilla@dloyavilla-virtual-machine: ~/Documents
xv6 kernel is booting
init: starting sh
$ time matmul
Time: 83
Time: 7 ticks
elapsed time: 7 ticks, cpu time: 7 ticks, 100% CPU
$ time1 matmul
Time: 716
Time: 7 ticks
Elapsed time: 7
time1 matmul
Time: 813
Time: 7 ticks
Elapsed time: 7
$ time1 matmul & ; time1 matmul &
Time: 959
Time: 960
$ Time: 13 ticks
Elapsed time: 14
Time: 12 ticks
Elapsed time: 13
$ time1 sleep 10
```

Figure 2. Test Cases

Task 2. Keep track of how much cputime a process has used.

Figure 5. Added int for cputime so we can store the cputime in proc.h

Figure 7. Initialized to 0 cpu time, in allocproc, in proc.c so each time is called we start the counter to 0

Figure 6 & 8. Incremented cputime in file trap.c, for user and kernel trap, so it increments everytime it goes into the user or kernel traps

By completing this job, I have gained knowledge on how to alter the xv6 operating system to monitor CPU usage by particular programs. The resource utilization of processes can be tracked and profiled using this.

My difficulty was how to increment the cputime was confusing but after changing the code a little I figured it out.

```
0 /* 280 */ uint64 t6;  
1};  
2  
3enum procstate { UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };  
4  
5// Per-process state  
6struct proc {  
7    struct spinlock lock;  
8  
9    // p->lock must be held when using these:  
10   enum procstate state; // Process state  
11   void *chan; // If non-zero, sleeping on chan  
12   int killed; // If non-zero, have been killed  
13   int xstate; // Exit status to be returned to parent's wait  
14   int pid; // Process ID  
15   int cputime; // cputime  
16  
17   // wait lock must be held when using this:  
18   struct proc *parent; // Parent process  
19  
20   // these are private to the process, so p->lock need not be held.  
21   uint64 kstack; // Virtual address of kernel stack  
22   uint64 sz; // Size of process memory (bytes)  
23   pagetable_t pagetable; // User page table  
24   struct tranframe *tranframe; // data base for tranframe.s  
25};
```

Figure 7. proc.h

```
5static struct proc*  
6allocproc(void)  
7{  
8    struct proc *p;  
9  
10   for(p = proc; p < &proc[NPROC]; p++) {  
11       acquire(&p->lock);  
12       if(p->state == UNUSED) {  
13           goto found;  
14       } else {  
15           release(&p->lock);  
16       }  
17   }  
18   return 0;  
19  
20 found:  
21   p->pid = allocpid();  
22   p->state = USED;  
23   p->cputime = 0;  
24 }
```

Figure 6. proc.c

```
pstat.h x time.h x user.h x sys.pl x syscall.c x syscall.h x sysproc.c x  
1// user trap handler  
2p->trapframe->sepc = r_sepc();  
3  
4if(r_scause() == 0){  
5    // system call  
6    if(p->killed)  
7        exit(-1);  
8  
9    // sepc points to the ecall instruction,  
10   // but we want to return to the next instruction.  
11   p->trapframe->sepc += 4;  
12  
13   // an interrupt will change sstatus & registers,  
14   // so don't enable until done with those registers.  
15   intr_on();  
16  
17   syscall();  
18 } else if((which_dev = devintr()) != 0){  
19     // ok  
20 } else {  
21     printf("usertrap(): unexpected scause %p pid=%d\n", r_scause(), p->pid);  
22     printf("sepc=%p stval=%p\n", r_sepc(), r_stval());  
23     p->killed = 1;  
24 }  
25  
26 if(p->killed)  
27     exit(-1);  
28  
29 // give up the CPU if this is a timer interrupt.  
30 if(which_dev == 2){  
31     yield();  
32     p->cputime++;  
33 }  
34 usertrapret();  
35 }
```

Figure 5. User trap

```
1// ON WHATEVER THE CURRENT KERNEL STACK IS.  
2void  
3kerneltrap()  
4{  
5    int which_dev = 0;  
6    uint64 sepc = r_sepc();  
7    uint64 sstatus = r_sstatus();  
8    uint64 scause = r_scause();  
9  
10   if((sstatus & SSTATUS_SPP) == 0)  
11       panic("kerneltrap: not from supervisor mode");  
12   if(intr_get() != 0)  
13       panic("kerneltrap: interrupts enabled");  
14  
15   if((which_dev = devintr()) == 0){  
16       printf("scause %p\n", scause);  
17       printf("sepc=%p stval=%p\n", r_sepc(), r_stval());  
18       panic("kerneltrap");  
19   }  
20  
21   // give up the CPU if this is a timer interrupt.  
22   if(which_dev == 2 && myproc() != 0 && myproc()->state == RUNNING){  
23       myproc()->cputime++;  
24       yield();  
25   }  
26  
27   // the yield() may have caused some traps to occur,  
28   // so restore trap registers for use by kernelvec.S's sepc instruction.  
29   w_sepc(sepc);  
30   w_sstatus(sstatus);  
31 }
```

Figure 8. Kernel trap

Task 3. Implement a wait2() system call that waits for a child to exit and returns the child's status and rusage.

Figure 10. added rusage structure in a new file called pstat.h

Figure 11. in user.h added structure so the user can use it.

Figure 12. usys.pl adding the identifier for the system call.

Figure 9. syscall.h, adding number corresponding to syscall

Figure 13. adding signature to make calls in user level.

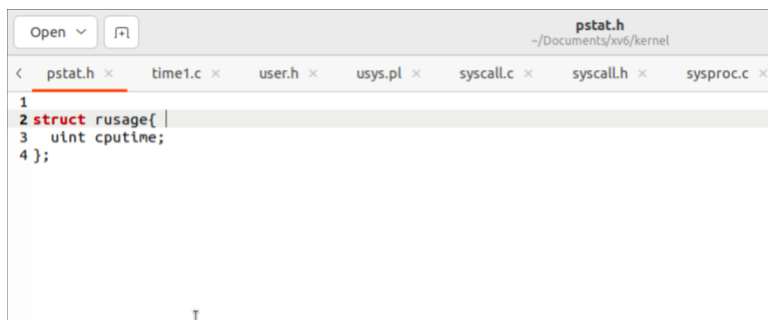
Figure 14. syscall.c this file is to request resources from kernel to user level, so we add wait2.

Figure 16. sysproc.c we added wait 2, and validated the arguments passed.

Figure 15. code to implement the new process, it is based on the existing wait, just adding the use of the structure rusage, and copying the status.

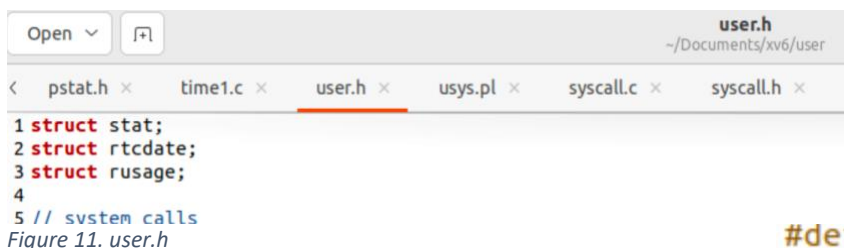
I learned about the functionality of each file, and how to code an alternative version of the wait system call, but this time also being able to calculate the cputime.

This was the most difficult task for this lab, because we had to understand the general system call infrastructure, mapping system call numbers, validating arguments, and ensuring that user-level and kernel-level code operate together. These difficulties can be overcome by carefully following the offered instructions, consulting documentation, and undertaking comprehensive testing to confirm that the implementation is right.



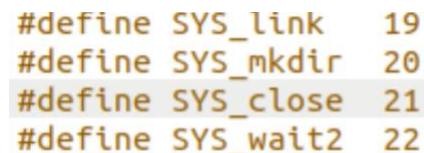
```
1
2 struct rusage {
3     uint cputime;
4 };
```

Figure 10. pstat.h



```
1 struct stat;
2 struct rtcdate;
3 struct rusage;
4
5 // system calls
```

Figure 11. user.h



```
#define SYS_link    19
#define SYS_mkdir   20
#define SYS_close   21
#define SYS_wait2   22
```

Figure 9. syscall.h

```

36 entry("sbrk");
37 entry("sleep");
38 entry("uptime");
39 entry("wait2");

```

Figure 12. usys.pl

```

4 char* sbrk(int);
5 int sleep(int);
6 int uptime(void);
7 int wait2(int*, struct rusage*);
8
9 // ulib.c
0 int stat(const char*, struct stat*);

```

Figure 13. user.h

```

extern uint64 sys_uptime(void);
extern uint64 sys_wait2(void);

static uint64 (*syscalls[])(void) = {
[SYS_fork]    sys_fork,
[SYS_exit]    sys_exit,
[SYS_wait]    sys_wait,
[SYS_pipe]    sys_pipe,
[SYS_read]    sys_read,
[SYS_kill]    sys_kill,
[SYS_exec]    sys_exec,
[SYS_fstat]    sys_fstat,
[SYS_chdir]    sys_chdir,
[SYS_dup]    sys_dup,
[SYS_getpid]    sys_getpid,
[SYS_sbrk]    sys_sbrk,
[SYS_sleep]    sys_sleep,
[SYS_uptime]    sys_uptime,
[SYS_open]    sys_open,
[SYS_write]    sys_write,
[SYS_mknod]    sys_mknod,
[SYS_unlink]    sys_unlink,
[SYS_link]    sys_link,
[SYS_mkdir]    sys_mkdir,
[SYS_close]    sys_close,
[SYS_wait2]    sys_wait2,
}

```

Figure 14. syscall.c

```

98
99 uint64
100 sys_wait2(void)
101 {
102     uint64 p1, p2;
103     if(argaddr(0, &p1) < 0 || argaddr(1, &p2) < 0){return -1;}
104     return wait2(p1,p2);
105 }

```

Figure 15. sysproc.c

```

431
432 //
433 int
434 wait2(uint64 addr, uint64 rusage)
435 {
436     struct proc *np;
437     int havekids, pid;
438     struct proc *p = myproc();
439     struct rusage ru;
440     acquire(&wait_lock);
441
442     for(;;){
443         // Scan through table looking for exited children.
444         havekids = 0;
445         for(np = proc; np < &proc[NPROC]; np++){
446
447             if(np->parent == p){
448                 // make sure the child isn't still in exit() or swtch().
449                 acquire(&np->lock);
450
451                 havekids = 1;
452                 if(np->state == ZOMBIE){
453
454                     ru.cputime = np->cputime;
455                     copyout(p->pagetable, rusage, (char *)&ru, sizeof(ru));
456                     //copying data from kernel mode to user mode
457
458                     // Found one.
459                     pid = np->pid;
460                     if(addr != 0 && copyout(p->pagetable, addr, (char *)&np->xstate,
461                                     sizeof(np->xstate)) < 0) {
462                         release(&np->lock);
463                         release(&wait_lock);
464                         return -1;
465                     }

```

Figure 16. proc.c

Task 4. Implement a time command that runs the command given to it as an argument and outputs elapsed time, CPU time, and %CPU used.

Figure 17. Code to implement the time, similar to time 1, we check if we have enough arguments, then shift the input, fork, check if the child was created successfully and if the exec was also successful.

Figure 18. added time to the makefile

Figure 19. testing the code with the given input

I learned how to create a program that waits for child processes to exit while managing processes, and measure the elapsed time, CPU time, and %CPU utilization is calculated and displayed. It was difficult to understand how processes are created, run, and waited for in the xv6 operating system.

```

1 #include "user/user.h"
2 #include "kernel/kstat.h"
3 #include "sysstat.h"
4 #include "sysproc.h"
5 #include "trap.h"
6 #include "def.h"
7 #include "cat.h"
8 #include "stat.h"
9 #include "time.h"
10
11 int t, start, elapsed;
12 struct rusage ru;
13
14 if (argc < 2) {
15     printf("usage: time <comm> [args...]\n");
16     exit(-1);
17 }
18
19 for (i = 1; i < argc; i++) {
20     newargv[i-1] = argv[i];
21 }
22 newargv[argc-i] = 0;
23 start = uptime();
24 printf("time: %d\n", start);
25 int pid = fork();
26 if (pid == -1) {
27     printf("fork failed\n");
28     exit(-1);
29 } else if (pid == 0) {
30     exec(newargv[0], newargv);
31     printf("exec failed\n");
32     exit(-1);
33 } else {
34     let upid = wait2(0, &ru);
35     elapsed = uptime() - start;
36     printf("elapsed time: %d ticks, cpu time: %d ticks, %d%% CPU\n", elapsed, ru.cputime, ru.cputime*100/elapsed);
37     if (upid == pid) {
38         // the shell exited; restart it.
39         exit(-1);
40     }
41 }

```

Figure 17. time.c

`$U/_sleep\`
`$U/_time\`

Figure 18. Makefile

```

$ matmul
Time: 6 ticks
$ time matmul
Time: 8999
Time: 7 ticks
elapsed time: 7 ticks, cpu time: 7 ticks, 100% CPU
$ time matmul & time matmul &
Time: 9185
Time: 9186
$ Time: 13 ticks
elapsed time: 14 ticks, cpu time: 7 ticks, 50% CPU
Time: 12 ticks
elapsed time: 13 ticks, cpu time: 7 ticks, 53% CPU

```

Figure 19. time test cases

Extra Credit (5 points). Discuss limitations of our time command. (Hint: For one limitation, consider what would happen if the command that is being timed forks child processes).

- Errors or inaccuracy with the cpu time because if we fork a process, the calculation will be only based on the parent, leaving all the child processes. Same with the rusage, it won't track each child separately.
- Also the code cannot be maintain if you want to measure complex commands, for example one that executes other commands.
- I'm not sure on how to solve this limitations but for sure we need to modify and make the code more complex, and I'm guessing with recursion we may be able to track each child process.