

### HW 3: Priority-based Scheduler for xv6

Task 1. Modify the provided ps command to print the priority of each process.

I added the "priority" field to the `struct proc` in `proc.h` to store process priorities and additional scheduling information.

In `sysproc.c`, I implemented the `getpriority()` and `setpriority()` system calls, allowing users to retrieve and set the priority of specific processes.

To make these system calls accessible to user-level programs, I added the necessary function prototypes and constants to `user.h` and updated `syscall.h`.

The changes were integrated into `syscall.c` to ensure that the new system calls are available for use by user-level programs.

```
// Per-process state
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum procstate state; // Process state
    void *chan;           // If non-zero, sleeping on chan
    int killed;           // If non-zero, have been killed
    int xstate;           // Exit status to be returned to parent's wait
    int pid;              // Process ID
    uint64 cputime;        // cputime
    int priority;          // add priority
    int readytime;

    // wait_lock must be held when using this:
    struct proc *parent;   // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;         // Virtual address of kernel stack
    uint64 sz;             // Size of process memory (bytes)
    pagetable_t pagetable; // User page table
    struct trapframe *trapframe; // data page for trampoline.S
    struct context context; // swtch() here to run process
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;       // Current directory
};
```

```
3 [SYS_close]    sys_close,
4 [SYS_getprocs] sys_getprocs,
5 [SYS_wait2]    sys_wait2,
6 [SYS_getpriority] sys_getpriority,
7 [SYS_setpriority] sys_setpriority,
8 };
9
```

```

28 int getprocs(struct pstat*);
29 int wait2(int*, struct rusage*);
30 int getpriority(void);
31 int setpriority(int);
32 |
33 // ulib.c

```

```

0
9
0 // getPriority()
1 uint64
2 sys_getpriority(void){
3     return myproc()->priority;
4 }
5
6 // setPriority()
7 uint64
8 sys_setpriority(void){
9     int priority;
0     if(argint(0, &priority) < 0) {
1         return -1;
2     }
3     myproc()->priority = priority;
4     return 0;
5 }
6
7 // return the number of active processes in the system

```

```

24 #define SYS_wait2 23
25 #define SYS_getpriority 24
26 #define SYS_setpriority 25

```

Task 2. Add a readytime field to struct proc, initialize it correctly, and modify ps to print a process's age.

I added the "readytime" field to the `struct proc` in `proc.h`. This field is initialized to the current time whenever a process's state transitions from another state to RUNNABLE.

In `kernel/pstat.h`, I modified the `struct pstat` to include the "readytime" field. This change enables the `ps` command to access and display process ages for processes in the RUNNABLE state.

The `ps` command is updated to print the process's age when its state is RUNNABLE. The process age is calculated by subtracting the "readytime" of the process from the current time. The result is displayed in seconds.

To find the age of the process, subtract its "ready time" from the current time when you need to display the age. This difference represents the time elapsed since the process became ready to run.

```
// Per-process State
struct proc {
    struct spinlock lock;

    // p->lock must be held when using these:
    enum procstate state; // Process state
    void *chan;           // If non-zero, sleeping on chan
    int killed;           // If non-zero, have been killed
    int xstate;           // Exit status to be returned to parent's wait
    int pid;              // Process ID
    uint64 cputime;        // cputime
    int priority;         // add priority
    int readytime;

    // wait_lock must be held when using this:
    struct proc *parent; // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;        // Virtual address of kernel stack
    uint64 sz;            // Size of process memory (bytes)
    pagetable_t pagetable; // User page table
    struct trapframe *trapframe; // data page for trampoline.s
    struct context context; // switch() here to run process
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
}
```

```
1 struct pstat {
2     int pid; // Process ID
3     enum procstate state; // Process state
4     uint64 size; // Size of process memory (bytes)
5     int ppid; // Parent process ID
6     int priority;
7 #include "kernel/param.h"
8 #include "kernel/types.h"
9 #include "kernel/pstat.h"
10 #include "user/user.h"
11
12 int
13 main(int argc, char **argv)
14 {
15     struct pstat uproc[NPROC];
16     int nprocs;
17     int i;
18     char *state;
19     static char *states[] = {
20         [SLEEPING] "sleeping",
21         [RUNNABLE] "runnable",
22         [RUNNING] "running",
23         [ZOMBIE] "zombie",
24     };
25     nprocs = getprocs(uproc);
26     if (nprocs < 0)
27         exit(-1);
28
29     printf("pid\tstate\tsize\tage\tpriority\tcputime\tppid\ttname\n");
30     for (i=0; i<nprocs; i++) {
31         int age = uptime() - uproc[i].readytime;
32         state = states[uproc[i].state];
33         printf("%d\t%s\t%d\t%d\t%d\t%d\t%d\t%s\n", uproc[i].pid, state,
34             uproc[i].size, age, uproc[i].priority, uproc[i].cputime, uproc[i].ppid, uproc[i].name);
35     }
36     exit(0);
37 }
```

### Task 3. Implement a priority-based scheduler.

Introduced constants in `param.h` to allow for the selection of scheduling policies at compile time. These constants determine the scheduling policy, such as whether the priority-based scheduler should be used.

Implemented a priority-based scheduler in the operating system. This scheduler selects the highest priority process for execution. In the case of a tie among the highest priority processes, the scheduler selects any one of them.

Modified `proc.h` and `proc.c` to include a "priority" field in the `struct proc` and initialized it for each process.

Modified the necessary system calls and their corresponding system call handlers to enable processes to set and get their priority values.

Developed test programs to validate the functionality of the priority-based scheduler. These programs include multiple processes with varying priorities.

The results of running the test programs indicate that the priority-based scheduler effectively prioritizes higher-priority processes, demonstrating the expected behavior of the scheduling policy. The use of the `pexec` program for testing helps showcase the priority-based execution of processes.

Difficulties: Implementing the scheduler needed extensive testing.

```
16 #define ROUNDROBIN    0
17 #define PRIORITY      1
18 #define DREFSCHED      PRIORITY
19 enum procstate {UNUSED, USED, SLEEPING, RUNNABLE, RUNNING, ZOMBIE};
```

```

496 void
497 scheduler(void)
498 {
499     struct proc *p;
500     struct cpu *c = mycpu();
501     struct proc *highProc;
502
503     c->proc = 0;
504     for(;;){
505         // Avoid deadlock by ensuring that devices can interrupt.
506         intr_on();
507         if(DREFSCHED == 0){
508             for(p = proc; p < &proc[NPROC]; p++) {
509                 acquire(&p->lock);
510                 if(p->state == RUNNABLE) {
511                     // Switch to chosen process. It is the process's job
512                     // to release its lock and then reacquire it
513                     // before jumping back to us.
514                     p->state = RUNNING;
515                     c->proc = p;
516                     swtch(&c->context, &p->context);
517
518                     // Process is done running for now.
519                     // It should have changed its p->state before coming back.
520                     c->proc = 0;
521                 }
522                 release(&p->lock);
523             }
524         } else {
525             highProc = proc;
526             int highestProcess = 0;
527             for(p = proc; p < &proc[NPROC]; p++) {
528                 int age = sys_uptime() - p->readytime;
529                 int priorityCal = p->priority + age;
530                 acquire(&p->lock);

```

```

516         swtch(&c->context, &p->context);
517
518         // Process is done running for now.
519         // It should have changed its p->state before coming back.
520         c->proc = 0;
521     }
522     release(&p->lock);
523 }
524 } else {
525     highProc = proc;
526     int highestProcess = 0;
527     for(p = proc; p < &proc[NPROC]; p++) {
528         int age = sys_uptime() - p->readytime;
529         int priorityCal = p->priority + age;
530         acquire(&p->lock);
531         if(p->state == RUNNABLE) {
532             if(priorityCal > highestProcess) {
533                 highestProcess = priorityCal;
534                 highProc = p;
535             }
536         }
537         release(&p->lock);
538     }
539     acquire(&highProc->lock);
540     if(highProc->state == RUNNABLE) {
541         highProc->state = RUNNING;
542         c->proc = highProc;
543         swtch(&c->context, &highProc->context);
544         c->proc = 0;
545     }
546     release(&highProc->lock);
547 }
548 }

```

```

init: starting sh
$ pexec 10 ps

```

pid	state	size	age	priority	cputime	ppid	name
1	sleeping		12288	49 0	0	0	init
2	sleeping		16384	9 0	0	1	sh
3	sleeping		12288	5 10	0	2	pexec
4	running		12288	1 0	0	3	ps

```

$

```

#### Task 4. Add aging to your priority based scheduler.

The priority-based scheduler is modified to include an aging policy to prevent process starvation.

The scheduler periodically checks the age of processes in the RUNNABLE state.

Processes in the RUNNABLE state have their priorities gradually increased over time to prevent starvation. The exact aging algorithm specifics may vary based on the operating system's design.

Developed test programs to verify the aging policy's functionality, assessing the impact of aging on process scheduling.

The results of running these test programs demonstrate the effectiveness of the aging policy.

Lower-priority processes are less likely to starve, leading to improved system performance and fair resource allocation.

Difficulties: Making the necessary adjustments to the code, so it can function correctly

```
496 void
497 scheduler(void)
498 {
499     struct proc *p;
500     struct cpu *c = mycpu();
501     struct proc *highProc;
502
503     c->proc = 0;
504     for(;;){
505         // Avoid deadlock by ensuring that devices can interrupt.
506         intr_on();
507         if(DREFSCHED == 0){
508             for(p = proc; p < &proc[NPROC]; p++) {
509                 acquire(&p->lock);
510                 if(p->state == RUNNABLE) {
511                     // Switch to chosen process. It is the process's job
512                     // to release its lock and then reacquire it
513                     // before jumping back to us.
514                     p->state = RUNNING;
515                     c->proc = p;
516                     swtch(&c->context, &p->context);
517
518                     // Process is done running for now.
519                     // It should have changed its p->state before coming back.
520                     c->proc = 0;
521                 }
522                 release(&p->lock);
523             }
524         } else {
525             highProc = proc;
526             int highestProcess = 0;
527             for(p = proc; p < &proc[NPROC]; p++) {
528                 int age = sys_uptime() - p->readytime;
529                 int priorityCal = p->priority + age;
530                 acquire(&p->lock);
```

```
516         swtch(&c->context, &p->context);
517
518         // Process is done running for now.
519         // It should have changed its p->state before coming back.
520         c->proc = 0;
521     }
522     release(&p->lock);
523 }
524 } else {
525     highProc = proc;
526     int highestProcess = 0;
527     for(p = proc; p < &proc[NPROC]; p++) {
528         int age = sys_uptime() - p->readytime;
529         int priorityCal = p->priority + age;
530         acquire(&p->lock);
531         if(p->state == RUNNABLE) {
532             if(priorityCal > highestProcess) {
533                 highestProcess = priorityCal;
534                 highProc = p;
535             }
536         }
537         release(&p->lock);
538     }
539     acquire(&highProc->lock);
540     if(highProc->state == RUNNABLE) {
541         highProc->state = RUNNING;
542         c->proc = highProc;
543         swtch(&c->context, &highProc->context);
544         c->proc = 0;
545     }
546     release(&highProc->lock);
547 }
548 }
```



```

1 #include "kernel/param.h"
2 #include "kernel/types.h"
3 #include "kernel/pstat.h"
4 #include "user/user.h"
5
6 int
7 main(int argc, char **argv)
8 {
9     struct pstat uproc[NPROC];
10    int nprocs;
11    int i;
12    char *state;
13    static char *states[] = {
14        [SLEEPING]  "sleeping",
15        [RUNNABLE]  "runnable",
16        [RUNNING]   "running ",
17        [ZOMBIE]    "zombie  ",
18    };
19
20    nprocs = getprocs(uproc);
21    if (nprocs < 0)
22        exit(-1);
23
24    printf("pid\tstate\tsize\tage\tpriority\tcputime\tppid\tname\n");
25    for (i=0; i<nprocs; i++) {
26        int age = uptime() - uproc[i].readytime;
27        state = states[uproc[i].state];
28        printf("%d\t%s\t%d\t%d\t%d\t%d\t%d\t\n", uproc[i].pid, state,
29              uproc[i].size, age, uproc[i].priority, uproc[i].cputime, uproc[i].ppid, uproc[i].name);
30    }
31    exit(0);
32 }
33
34

```

xv6 kernel is booting

init: starting sh

\$ pexec 5 matmul 5 &; matmul 10 &

\$ pexec 10 ps

pid	state	size	ppid	name	priority	cputime	age
1	sleeping	12288	0	init	0		
2	sleeping	16384	1	sh	0		
7	runnable	12288	5	matmul	0		
6	runnable	12288	1	matmul	0		
5	sleeping	12288	1	pexec	0		
8	sleeping	12288	2	pexec	0		
9	running	12288	8	ps	0	0	177

\$ Time: 88 ticks

Time: 184 ticks

pexec 10 ps

exec pexec failed

\$ pexec 10 ps

pid	state	size	ppid	name	priority	cputime	age
1	sleeping	12288	0	init	0		
2	sleeping	16384	1	sh	0		
11	sleeping	12288	2	pexec	0		
12	running	12288	11	ps	0	0	557

\$